

Robust parsing techniques for semantic analysis of natural language queries

Afzal Ballim & Vincenzo Pallotta

MEDIA research group

Theoretical Computer Science Laboratory (LITH)

Swiss Federal Institute of Technology Lausanne (EPFL)

{Ballim,Pallotta}@di.epfl.ch

Abstract

This paper presents the result of an experimental system aimed at performing a robust semantic analysis of analyzed speech input in the area of information system access. The goal of this experiment was to investigate the effectiveness of such a system in a pipelined architecture, where no control is possible over the morpho-syntactic analysis which precedes the semantic analysis and query formation. The approach taken used technology from robust syntactic parsing applied to a sequence of parse trees rather than a sequence of lexical items, together with a constraint-based logical inference system that evaluated logical hypotheses about the query using fuzzy combination criteria.

1. Introduction

This paper describes how robust parsing techniques can be fruitfully applied for building a query generation module which is part of a pipelined NLP architecture aiming to process natural language queries in a restricted domain. We want to show that robustness represents a key issue in those NLP systems where it is more likely to have partial and ill-formed utterances due to various factors (e.g. noisy environments, low quality of speech recognition modules, etc...) and where it is necessary to succeed, even if partially, in extracting some meaningful information.

The domain we concerned in our case study is the interaction through speech with information systems. The availability of a huge collection of annotated telephone calls for querying the Swiss phone-book database (i.e. the Swiss French PolyPhone corpus [9]) allowed us to experiment our recent findings in robust text analysis obtained in the context of the Swiss National Fund research project ROTA (Robust Text Analysis), and in the Swisscom funded project ISIS (Interaction through Speech with Information Systems). Within this domain, the goal is to build a valid query to an information system, using limited world knowledge of the domain in question. Although a task like this may, at its simplest, be performed quite effectively using heuristic methods such as keyword spotting, such an approach is brittle, and does not scale up easily in the case of conducting a dialogue.

Problem specification

In this section we will give an informal specification for the problem of processing telephone calls for querying a phone-book database.

Swiss French PolyPhone Database

This database contains 4293 simulated recordings related to the "111" Swisscom service calls (e.g. "rubrique 38" of the calling sheet [9]). Each recording consists of 2 files, one ASCII text file corresponding to the initial prompt and the information request and one data file containing the sampled sound version. As far as the address fields are concerned, the data in the PolyPhone database are unfortunately not tagged and even not consistent. Prompts and information requests expressed by users have been extracted from the files and regrouped into a single representation in the following format:

```
id:cd1/b00/f0000o06:sid17733 prompt:11
adr1:MOTTAZ MONIQUE
adr2:rue du PRINTEMPS 4
adr3:SAIGNELEGIER
text[123]: Bonjour j'aimerais un numéro de
téléphone à Saignelegier c'est Mottaz m o deux ta z
Monique rue du printemps numéro quatre
sample:0.200000:10.820000:88160:42801
```

where currently, the corresponding lines in text file are processed with the following heuristic:

id identifies the original location of the file in the CD-ROM.

prompt identifies both the type of prompt asking the user for posing the query (e.g. n. 1 corresponds to "Veuillez maintenant faire comme si vous étiez en ligne avec le 111 pour demander le numéro de téléphone de la personne imaginaire dont les coordonnées se trouvent ci-dessous:").

adr1 corresponds to the name

adr2 corresponds to the address if line 3 is not empty and town otherwise

adr3 corresponds to the town if not empty. **text** corresponds to the text transcription. The number in square brackets is the total number of chars in the request.

sample groups the information for the sampled sound version of the request.

This heuristic seems to perform quite well but a more thorough and exhaustive evaluation still needs to be carried out. The main problem remains in finding enough information about the original data in order to be able to perform the validation automatically.

Concerning the structure in the Swiss Phone-book database, we assumed it is the same as the one that appears on the web (e.g. <http://www.ife.ee.ethz.ch/cgi-bin/etvq/f>), namely (one field per line):

Nom de famille / Firme
Prénom / Autres informations
No de téléphone
Rue, numéro
NPA, localité

One point still remains unclear about the PolyPhone database (as no answers were found in [3, 9]): what was the set of annotation used for the transcription of utterances? Several speech annotations such as "<hesitation>" appear in the text. Was it systematic? Are there other such markers? It is possible to rely on prosodic informations? In the first phase of the project we simply skipped these informations but we guess that they could be of great help in disambiguating interpretations of strict adjacent sequences of names such as in utterances like "j'amerais le numéro de téléphone de Vedo-Moser Brigitte Brignon Baar-Nendaz".

Overall architecture

The processing of the corpus data is performed at various linguistic levels performed by modules organized into a pipeline. Each module assumes as input the output of the preceding module. The main goal of this architecture is to understand how far it is possible to go without using any kind of feedback and interactions among different linguistic modules. In the case study we are going to present, only three linguistic levels are considered, namely the morpho-syntactic, the syntactic and the semantic level.

Morpho-Syntactic analysis

Morphological and syntactic processing were performed by our partners in the ISIS project, and the results are reported separately. In brief, low-level processing (morphological analysis and tagging) were performed by ISSCO (Institute Dalle Molle, University of Geneva) using tools that were developed in the European Linguistics Engineering project MULTEXT. For syntactic analysis, ISSCO developed a Feature Unification Grammar based on a small sample of the Polyphone data. This

grammar was taken by another of our partners (the Laboratory for Artificial Intelligence of the Swiss Federal Institute of Technology, Lausanne) and converted into a probabilistic context-free grammar, which was then applied to a sample of 500 entries from the Polyphone data.

Syntactic analysis

The forest of syntactic trees produced by this phase served as the input for the robust semantic analysis that we performed, that had as goal the production of query frames for the information system. Domain knowledge was employed to ensure that the query frames contained a minimal amount of information necessary for the query to be considered well-formed. This same world knowledge could further be used to identify limited cases of queries that were out of the domain (for instance, asking about a phone number in France from an information system that only has data concerning numbers in Switzerland).

Semantic annotations

The goal of semantic annotations is to provide as much information regarding the meaning of phrases as could reasonably be assumed possible if an analysis were performed using:

- state of the art techniques in semantic analysis;
- a semantic model of the closed world presented by the domain of the information system to be queried. This model would be used, for instance, to constrain the senses of words allowable in context, thus reducing the ambiguity of the phrase;
- a guessing module which could provide semantic hints regarding words that are not found in the lexicon (e.g., proper name identification, verbalized pauses and other non-word articulations which may be used in the phrasal corrective phase of robust analysis);
- a robust analysis module which can extract a coherent interpretation of the utterance.

While the semantic analysis will in general reduce the degree of ambiguity found after syntactic analysis, there remains the possibility that it might increase some degree of ambiguity due to the presence of coherent senses of words with the same syntactic category (e.g., the word "Geneva" can refer to either the canton or the city). In cases where this ambiguity has little effect on the overall interpretation of the phrase (e.g., the user wishing to have a number to obtain the weather forecast for a region such as Geneva) a compact representation can be used, with the ambiguity represented by a disjunction over the senses of Geneva.

It should be noted that the effect of such ambiguity must be taken in a dialogue context. Thus, the ambiguity between the interpretation of Geneva as a city or a canton might be insignificant in the context of a finding the phone number for a

regional weather forecast system, and such regional indistinctness could be part of the semantic domain model (presuming an domain model which takes account of special information phone numbers). Thus, in such a case the semantic module could decide to eliminate the ambiguity as not being relevant to the ultimate querying goal of the system. On the other hand, the difference between the canton and city of Bern might be more significant, and therefore merit being explicitly represented.

Semantic robust analysis and frame filling

Although the initial annotation process has eliminated much of the repetition, restarts, verbalized pauses, and other characteristics that characterize speech, any real system must be able to deal with these aspects. The component that deals with such input is generally referred to as a robust analyzer. Although robustness can be considered as being applied at either a syntactic or semantic level, it is generally at the semantic level that it is most effective (although, for reasons of efficiency of the syntactic analysis, some robust analysis is needed to demark discourse segments for separate syntactic analysis).

This robust analysis needs a model of the domain in which the system operates, and a way of linking this model to the lexicon used by the other components. It specifies semantic constraints that apply in the world and which allow us to rule out incoherent requests (for instance). The degree of detail required of the domain model used by the robust analyzer depends upon the ultimate task that must be performed -- in our case, furnishing a query to an information system. Taking the assumption that the information system being queried is relatively close in form to a relational database, the goal of the interpretative process is to furnish a query to the information system that can be viewed in the form of a frame with certain fields completed, to function of the querying engine being to fill in the empty fields.

One way in which the interface could interact with the querying system would be to submit such a frame at the end of the analysis process without performing any coherency checking. The advantage of this method is that the model of the domain of queries that is required by the interface can be limited. However, such an approach has two major disadvantages:

1. the result of incorrectly formulated queries may be completely uninterpretable or erroneous, and the interface system would have no basis for evaluating the quality of such replies, or how to aid the user in formulating a better one;
2. there might be a number of possible frames that could be submitted for any instance of a user utterance/query, and this number might be

reducible by application of a model of coherent queries.

We will, therefore, presume that queries must be classified by the interface into three categories:

1. the query is correct -- the fields of the frame which must be completed contain semantically valid data. The query may be submitted;
2. incomplete queries -- certain necessary fields cannot be unambiguously filled in, and so a system-initiative dialogue can be invoked to furnish the necessary information to create a correct query;
3. incoherent queries -- information in the fields of the frame is not coherent with the interfaces model of the domain. An error dialogue must be invoked.

The last query category is the most complex, since it requires a domain model sufficiently rich to decide whether a query is outside of the domain, or inside the domain but violating certain semantic constraints. In addition, it requires relatively complex dialogue management as the corrective dialogue may involve resolution of miscomprehension by either the system or the user.

2. Computational logic for robust analysis

Historically the role of computational logic in computer science has been that of declarative languages to write executable specification. In computational linguistic all the syntactic formalism can be considered as deductive systems (see [6]) and logic formalism are often used to express semantics and pragmatics. What has been considered to be an advantage using logic -based programming languages is the symbol processing capability and the way of abstracting from the actual implementation of needed data structures. This perspective attracted the linguists who were not skilled programmer and who needed to rapidly represent and check their linguistic theories. The main drawback to this approach is efficiency but it is not the only one. In recent years several efforts have been done to improve efficiency of logic and functional programming languages by means of powerful abstract machines and optimized compilers. Sometimes, efficiency recovery leads to introduction of non-logical features in the language and the programmer should be aware of it in order to exploit it in the development of his or her applications (i.e. cut in logic programming).

Logic Programming and NLP

Definite Clause Grammars come up to mind when relating Logic Programming and Natural Language Processing. This is of course one of the best coupling between Computational Linguistics and Logic, but considered from the perspective of Robust NLP this union reveals some deficiencies. In 1995, Erbach and Manandhar [10] examined the

state of the art in Logic Programming techniques for NLP and a wish-list of features for future developments were proposed. Nowadays some of those desired features are available in actual systems and can be considered referring to the two following activities:

1. Support for the development of linguistic models of Natural Language (Computational Linguistics)
2. Support for the design of real life applications (Language Engineering)

Declarativeness gives an undoubted benefit to both activities, but from the perspective of Robust Analysis, it allows to specify the robustness problem in a more rational way. Rather than concerning with implementation details, robustness can be achieved just suitably composing (possibly concurrently) logical modules (theories) which have a clear mathematical semantics. Thus robustness can be achieved by means of a cooperating linguistic architecture and suitably stated by flexible linguistic formalism.

Left-corner Head-driven Island Parser

In this section we overview the key features of the LHIP system [5, 15] which has been used as the main tool for doing semantical analysis in the context of the project ISIS.

LHIP is a system which performs robust analysis of its input, using a grammar defined in an extended form of the Definite Clause Grammar formalism used for implementation of parsers in Prolog. The chief modifications to the standard Prolog 'grammar rule' format are of two types: one or more right-hand side (RHS) items may be marked as 'heads', and one or more RHS items may be marked as 'ignorable'. We expand on these points and introduce other differences below.

LHIP employs a different control strategy from that used by Prolog DCGs, in order to allow it to cope with ungrammatical or unforeseen input. The behavior of LHIP can best be understood in terms of the complementary notions of span and cover. A grammar rule is said to produce an island which spans input terminals t_i to t_{i+n} if the island starts at the i^{th} terminal, and the $i + n^{th}$ terminal is the terminal immediately to the right of the last terminal of the island. A rule is said to cover m items if m terminals are consumed in the span of the rule. Thus $m \leq n$. If $m = n$ then the rule has completely covered the span.

As implied here, rules need not cover all of the input in order to succeed. More specifically, the constraints applied in creating islands are such that islands do not have to be adjacent, but may be separated by noncovered input. There are two notions of non-coverage of the input: unsanctioned and sanctioned noncoverage. The former case arises when the grammar simply does not account for some terminal. Sanctioned non-coverage means

that special rules, called "ignore" rules, have been applied so that by ignoring parts of the input the islands are adjacent. Those parts of the input that have been ignored are considered to have been consumed. These ignore rules can be invoked individually or as a class. It is this latter capability which distinguishes ignore rules from regular rules, as they are functionally equivalent otherwise, but mainly serve as a notational aid for the grammar writer.

Strict adjacency between RHS clauses can be specified in the grammar. It is possible to define global and local thresholds for the proportion of the spanned input that must be covered by rules; in this way, the user of an LHIP grammar can exercise quite fine control over the required accuracy and completeness of the analysis.

A chart is kept of successes and failures of rules, both to improve efficiency and provide a means of identifying unattached constituents. In addition, feedback is given to the grammar writer on the degree to which the grammar is able to cope with the given input; in a context of grammar development, this may serve as notification of areas to which the coverage of the grammar might next be extended.

Extensions of Prolog DCG grammars in LHIP permit:¹

1. nominating certain RHS clauses as heads;
2. marking some RHS clauses as being optional;
3. invocation of ignore rules;
4. imposing adjacency constraints between two RHS clauses; and
5. setting a local threshold level in a rule for the fraction of spanned input that must be covered. A threshold defines the minimum fraction of terminals covered by the rule in relation to the terminals spanned by the rule in order for the rule to succeed. For instance, if a rule spans terminals t_i to t_{i+n} covering j terminals in that span, then the rule can only succeed if $j/n \geq T$.

The following is an example of a LHIP rule. At first sight this rule appears left recursive. However, the sub-rule "conjunction(Conj)" is marked as a head and therefore is evaluated before either of "s(SI)" or "s(Sr)". Presuming that the conjunction-rule does not end up invoking (directly or indirectly) the s-rule, then the s-rule is not left-recursive. On the other hand, if the conjunction-rule were not marked as a head, then the s-rule would be left recursive and would not terminate.

```
s(conjunct(Conj,SI,Sr)) ~>
    s(SI),
    *conjunction(Conj),
    s(Sr).
```

¹A version of LHIP exists which permits a form of negation on RHS clauses. That version is not described here.

A number of tools are provided for producing analyses of input by the grammar with certain constraints. For example, to find the set of analyses that provide maximal coverage over the input, to find the subset of the maximal coverage set that have minimum spans, and to find the analyses that have maximal thresholds. In addition, other tools can be used to search the chart for constituents that have been found but are not attached to any complete analysis. The conversion of the grammar into Prolog code means that the user of the system can easily develop analysis tools that apply different constraints, using the given tools as building blocks.

Tree-paths representation Parse trees obtained from the previous module are encoded into a path representation which allows us to easily specify constraints over the tree structure. A path-sentence is a list of path-words which in turn are compound

```

graph TD
    Root(( )) --- 60M-11[60M-11]
    Root --- 25M-11[25M-11]
    25M-11 --- 5M-11L[5M-11]
    25M-11 --- 5M-11R[5M-11]
    5M-11L --- N-malcom[N-malcom]
    5M-11L --- 28-25M-11L[28-25M-11]
    5M-11R --- 28-25M-11R[28-25M-11]
    5M-11R --- 10M-11[10M-11]
  
```

Using this representation it is possible to define a grouping operator which given a sequence of adjacent names finds the subsequence of words having the least common ancestor which is closer than the least common ancestor of the given sequence. The group operator and least common ancestor predicates are very useful for imposing structural knowledge constraints and they are straightforwardly defined as PROLOG programs by:

```

prefix_path(A,A).
prefix_path([_|B],C):-
    prefix_path(B,C).

```

21

definition, we intend to follow the experiments done in the context of Verbmobil (see for example [12, 13]). It is worth noting that identification of such discourse segments by a robust syntactic analysis pre-processing system, can allow for more efficient syntactic analysis, since specific grammars can be written to deal with the different segments. This can provide significant speed ups, since speed is dependent on the cube of the number of grammar rules.

In our specific case identifying special words serving both as separators among logical subparts of the same sentence and as introducers of semantic constituents allows us to search for name sequences to fill a particular slot only in interesting part of the sentence. One of the most important separator is the announcement-query separator. The LHIP clauses defining this separator can be one or more word covering rule like for instance:

```
ann_query_separator #1.0 ~>
  @terminal('te'le'phone',_).
```

```
ann_query_separator #1.0 ~>
  (@terminal('nume'ro',_):
   @terminal('de',_):
  (? @terminal('te'le'phone',_) ?)).
```

As an example of semantic constituents introducers we propose here the

```
street_intro([T,Prep,Det],1)#1.0
~>
  *street_type(T),
  preposition(Prep),
  determiner(Det).
```

which make use of some word knowledge about street types coming from an external thesaurus like:

```
street_type(terminal(X,P)) ~>
  @terminal(X,P), {thesaurus(street,W),
  member(X,W)}.
```

Generation of hypotheses The generation of hypotheses for filling the frame is performed by:

- composing weighted rules
- assembling chunks
- filtering possible hypotheses

Weighted rules

The main assumption on which probabilistic approach to NLP is based, is that language is considered as being a random phenomenon with its own probability distribution function. Thus coverage of linguistic phenomena is often translated as expectation of those phenomena in a probabilistic sense. Changing perspective and considering language just as an uncertain and imprecise phenomenon and understanding as a perception process, it is naturally to think of

"fuzzy" models of language. Fuzzy set theory and hence Fuzzy Logic applied to language processing (see [14]) seems to be a promising approach (and has already been investigated in [4]). Recently, fuzzy reasoning has been partially integrated into a CLP paradigm (see [17]) in order to deal with so called soft constraints in weighted constraint logic grammars. We tried to get some inspiration from the above proposal for integrating fuzzy logic and parsing to compute weights to assign to each frame filling hypotheses. Each LHIP rule returns a confidence factor together with the sequence of names. The confidence factor for a rule can be assigned statically to pre-terminal rules (e.g. those identifying separators or introducers) or can be computed composing recursively the confidence factors of sub-constituents. Confidence factors are combined choosing the minimum among confidences of each sub-constituents. It is possible that there is no enough information for filling a slot. In this case the grammar should provide a mean to provide an empty constituent when all possible hypothesis rules have failed. This is possible using negation and epsilon-rules in LHIP as showed in the following rules for dealing with street names.

```
found_street_name(L,Conf) #1.0 ~>
  *street_intro(Intro,Conf),
  name_list(X),
  {append(Intro,X,L)}.
```

```
found_street_name(X,0.3) ~>
  *name_list(X).
```

```
hyp_street_name(Street,Conf) ~>
  *found_street_name(Street,Conf).
```

```
hyp_street_name([],1) ~>
  ~found_street_name(_,_),
  lhip_true.
```

where `name_list(X)` accounts for sequence of adjacent proper names and `lhip_true` corresponds to the empty sequence.

Observe that in this particular case there is no need to select the minimum confidence factor from the subconstituents of the rule `found_street_name` since we have only `street_intro(Intro,Conf)` which propagates its confidence factor.

Chunk assembling

The highest level constituent is represented by the whole frame structure which simply specifies the possible orders of chunks relative to slot hypotheses. The main rule for a possible frame hypothesis is:

```
frame( Caller_title,
        Caller_name,
        Target_title,
        Target_name,
        Street_name,
        Street_number,
        Locality,Weight) ~>
```

```

hyp_caller(Caller_title,Caller_name,C1),
*ann_query_separator,
hyp_target(Target_title,Target_name,C2),
*location_intro,
((hyp_street_name(Street_name,C3),
hyp_street_number(Street_number,C4));
(hyp_street_number(Street_number,C4),hyp_street_n
ame(Street_name,C3))),
hyp_locality_name(Locality,C5),

{lca(Target_name,P2),
length(Target_name,N2),
length(P2,Span2),
W2 = weight(C2,(1-(1/Span2))),(1-(1/N2))),
lca(Street_name,P3),
length(Street_name,N3),
length(P3,Span3), W3 =
weight(C3,(1-(1/Span3))),(1-(1/N3))),
minlist([weight(C1,0,0),W2,W3, weight(C4,0,0),
weight(C5,0,0), ],Weight)}.

```

In this rule we specify a possible order of chunks interleaved by separators and introducers. It is possible to include a disjunction of possible sub-orders using the ";" connective. The computation of global weight for the frame rule is more complex than simple chunk rules. In this case we consider for certain chunks also structural information (e.g. preferring names chunks belonging to the minimal common sub-tree) and we try to find the longest sequence of name belonging to the same sub-tree.

Filtering

The obtained frame hypotheses can be further filtered by using

- structural knowledge, that is constraints over the tree-path representation
- word knowledge.

Query generation

We now outline how to combine the information extracted from the previous analysis step into the final query representation which can be directly mapped into the database query language. We will make use of a frame structure in which slots represent information units or attributes in the database. A simple notion of context can be useful to fill by default those slots for which we have no explicit information. For instance it is possible to fill automatically the locality slot using information about the caller's prefix number in case of he does not explicitly specify any locality. We assume that his request actually refers to the same locality denoted by the above prefix number. For doing this type of hierarchical reasoning we exploit the meta-programming capabilities of logic programming and we used a meta-interpreter which allows multiple inheritance among logical theories [8].

4. Conclusions

From a very superficial observation of the human language understanding process, it appears clear that no deep competence of the underlying structure of the spoken language is required in order to be able to process acceptably distorted utterances. On the other hand, the more experienced is the speaker, the more probable is a successful understanding of that distorted input. How can this kind of fault-tolerant behavior be reproduced in an artificial system by means of computational techniques? Several answers have been proposed to this question and many systems implemented so far, but no one of them is capable of dealing with robustness as a whole.

Although the number of NLP-papers referring to robustness is quite large, an agreement upon its meaning is still missing. A quite reasonable definition is:

[Robustness is] . . . a kind of monotonic behavior, which should be guaranteed whenever a system is exposed to some sort of non-standard input data: A comparatively small deviation from a predefined ideal should lead to no or only minor disturbances in the system's response, whereas a total failure might only be accepted for sufficiently distorted input [16].

As examples of robust approaches applied to dialogue systems we cite here two systems which are based on similar principles.

In the DIALOGOS human-machine telephone system (see [1]) the robust behavior of the dialogue management module is based both on a contextual knowledge base of pragmatic-based expectations and the dialogue history. The system identifies discrepancies between expectations and the actual user behavior and in that case it tries to rebuild the dialogue consistency. Since both the domain of discourse and the user's goals (e.g. railway timetable inquiry) are clear, it is assumed the systems and the users cooperate in achieving reciprocal understanding. Under this underlying assumption the system pro-actively asks for the query parameters and it is able to account for those spontaneously proposed by the user.

In the SYSLID project (see [7]) where a robust parser constitutes the linguistic component (LC) of the queryanswering dialogue system. An utterance is analyzed while at the same time its semantical representation is constructed. This semantical representation is further analyzed by the dialogue control module (DC) which then builds the database query. Starting from a word graph generated by the speech recognizer module, the robust parser will produce a search path into the word graph. If no complete path can be found, the robust component of the parser, which is an island based chart parser (see [11]), will select the maximal consistent partial results. In this case the

parsing process is also guided by a lexical semantic knowledge base component that helps the parse in solving structural ambiguities.

We can conclude that robustness in dialogue is crucial when the artificial system takes part in the interaction since inability or low performance in processing utterances will cause unacceptable degradation of the overall system. As pointed out in [2] it is better to have a dialogue system that tries to guess a specific interpretation in case of ambiguity rather than ask the user for a clarification. If this first commitment results later to be a mistake a robust behavior will be able to interpret subsequent corrections as repair procedures to be issued in order to get the intended interpretation.

References

- [1] Dario Albesano, Paolo Baggia, Morena Danieli, Roberto Gemello, Elisabetta Gerbino, and Claudio Rullent. Dialogos: a robust system for human-machine spoken dialogue on the telephone. In Proc. of ICASSP, Munich, Germany, 1997.
- [2] J.F. Allen, B. Miller, E. Ringger, and T. Sikorski. A robust system for natural spoken dialogue. In Proc. 34th Meeting of the Assoc. for Computational Linguistics. Association of Computational Linguistics, June 1996.
- [3] J.M. Andersen, G. Caloz, and H. Boulard. Swisscom "advanced vocal interfaces services" project. Technical Report COM-97-06, IDIAP, Martigny, December 1997.
- [4] Peter. R.J. Asveld. Towards robustness in parsing - fuzzifying context-free language recognition. In J. Dassow, G. Rozemberg, and A. Salomaa, editors, *Developments in Language Theory II - At the Crossroad of Mathematics, Computer Science and Biology*, pages 443-453. World Scientific, Singapore, 1996.
- [5] A. Ballim and G. Russell. LHIP: Extended DCGs for Configurable Robust Parsing. In Proceedings of the 15th International Conference on Computational Linguistics, pages 501 - 507, Kyoto, Japan, 1994. ACL.
- [6] Patrick Blackburn, Marc Dymetman, Alain Lecomte, Aarne Ranta, Christian Retore', and Eric Villemonte de la Clergerie. Logical aspects of computational linguistics: an introduction. In Christian Re'tore', editor, *Logical Aspects of Computational Linguistics*, volume 1328 of LNCS/LNAI, pages 1 - 20. Springer, 1997.
- [7] Manuela Boros, Gerhard Hanrieder, and Ulla Ackermann. Linguistic processing for spoken dialogue systems - experiences made in the syslid project -. In Proceedings of the third CRIM-FORWISS Workshop, Montreal, Canada, 1996.
- [8] A. Brogi and F. Turini. Meta-logic for program composition: Semantic issues. In K.R. Apt and F.Turini, editors, *Meta-Logics and Logic Programming*. The MIT Press, 1995.
- [9] G. Chollet, J.-L. Chochard, A. Constantinescu, C. Jaboulet, and Ph. Langlais. Swiss french polyphone and polyvar: Telephone speech database to model inter- and intra-speaker variability. Technical Report RR-96-01, IDIAP, Martigny, April 1996.
- [10] G.Erbach and S.Manandhar. Visions for logic-based natural language processing. In Proceedings of the ILPS '95 workshop: "Visions for the Future of Logic Programming - Laying the Foundations for a Modern Successor to Prolog", Portland, Oregon, 1995.
- [11] G. Hanrieder and G Go"rz. Robust parsing of spoken dialogue using contextual knowledge and recognition probabilities. In Proceedings of the ESCA Tutorial and Research Workshop on Spoken Dialogue Systems - Theories and Applications, pages 57-60, Denmark, May 1995.
- [12] S. Jekat, A. Klein, E. Maier, I. Maleck, M. Mast, and J.J. Quantz. Dialogue acts in vermobil. *Vermobil Report* 65, DFKI, 1995.
- [13] R. Kompe, A. Kiebling, T. Kuhn, M. Mast, H. Niemann, E. No"th, K. Ott, and A. Batliner. Prosody takes over: A prosodically guided dialog system. *Vermobil report* 47, DFKI, 1994.
- [14] E.T. Lee and L.A. Zadeh. Note on fuzzy languages. *Information Science*, 1:421-434, 1969. [15] C. Lieske and A. Ballim. Rethinking natural language processing with prolog. In Proceedings of Practical Applications of Prolog and Practical Applications of Constraint Technology (PAPACTS98), London,UK, 1998. Practical Application Company.
- [16] Wolfgang Menzel. Robust Processing of Natural Language. In Ipke Wachsmuth, Claus-Rainer Rollinger, and Wilfried Bauer, editors, *KI-95: Advances in Artificial Intelligence*, Berlin, 1995. Springer.
- [17] Stefan Riezler. Quantitative constraint logic programming for weighted grammar applications. In *Logical Aspects of Computational Linguistics (LACL'96)*, LNCS. Springer, 1996.